

learn network inspire

GameDevelopers
Conference 08



February 18-22, 2008
San Francisco

www.gdconf.com



Advanced Soft Shadow Mapping Techniques

Louis Bavoil
NVIDIA Developer Technology



Why Soft Shadows?

- ⦿ Antialiasing

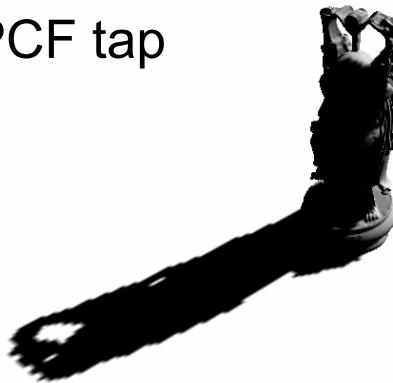
- Filtering the shadow edges

- ⦿ Area lights

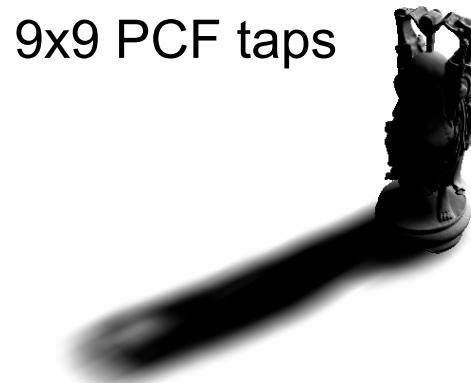
- Cast penumbra with variable size

- Shadows hardening on contact

1 PCF tap



9x9 PCF taps



PCSS





Why Shadow Mapping for Soft Shadows?

- ➊ Most popular technique for shadows in games
 - ➌ Purely image-based technique
 - ➌ Works with any rasterizable geometry

- ➋ Alternative techniques exist
 - ➌ Silhouette based techniques such as smoothies and penumbra wedges
 - ➌ Silhouette detection robustness issues
 - ➌ Do not work with alpha-tested geometry



Outline

➊ Fixed-Size Penumbra

- ➌ PCF (Percentage Closer Filtering)
- ➌ VSM (Variance Shadow Maps)
- ➌ CSM (Convolution Shadow Maps)
- ➌ ESM (Exponential Shadow Maps)

➋ Variable-Size Penumbra

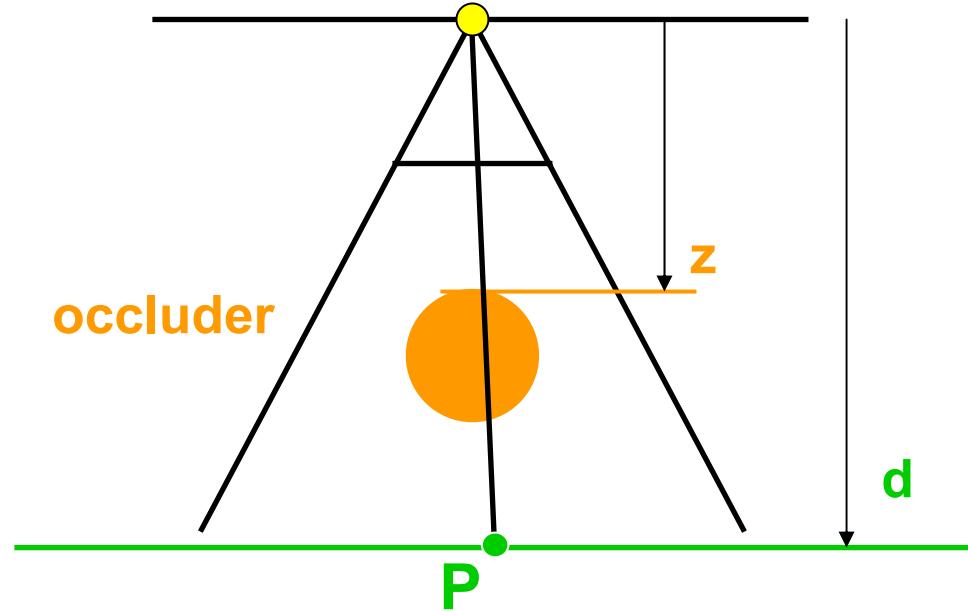
- ➌ PCSS (Percentage Closer Soft Shadows)
- ➌ PCSS + VSM/CSM
- ➌ Backprojection



Shadow Mapping

Game Developers' Conference 08

- Shadow map stores distance **z** to the light
 - P is lit $\Leftrightarrow (\mathbf{d} < \mathbf{z})$





Percentage Closer Filtering (PCF)

- ➊ Sample the result of ($d < z$) around projected point
 - ➋ Filter the binary results in a given kernel

0	1
1	0

- ➌ Bilinear PCF
 - ➍ NVIDIA and recent AMD GPUs implement 2x2 PCF in one fetch
 - ➎ Using same sample locations and weights as for bilinear filtering



GameDevelopers
Conference 08

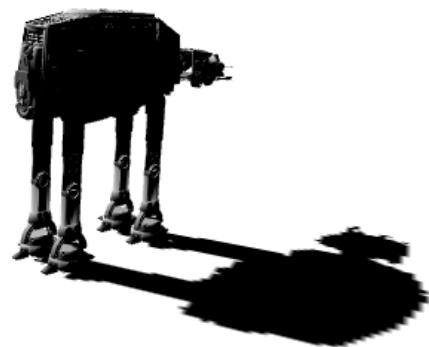
Using Bilinear PCF with DX10

```
Texture2D<float> tDepthMap;  
  
SamplerComparisonState ShadowSampler  
{  
    ComparisonFunc = LESS;  
    Filter = COMPARISON_MIN_MAG_LINEAR_MIP_POINT;  
};  
  
// ...  
sum += tDepthMap.SampleCmpLevelZero(ShadowSampler,  
                                    uv + offset, z);
```



PCF Filtering

- Increasing the number of PCF taps increases the softness of the shadows



1 tap



9x9 taps



17x17 taps



Irregular PCF

- ⦿ PCF with large kernels requires many samples
- ⦿ Using irregular sampling
Trades banding for noise



regular sampling

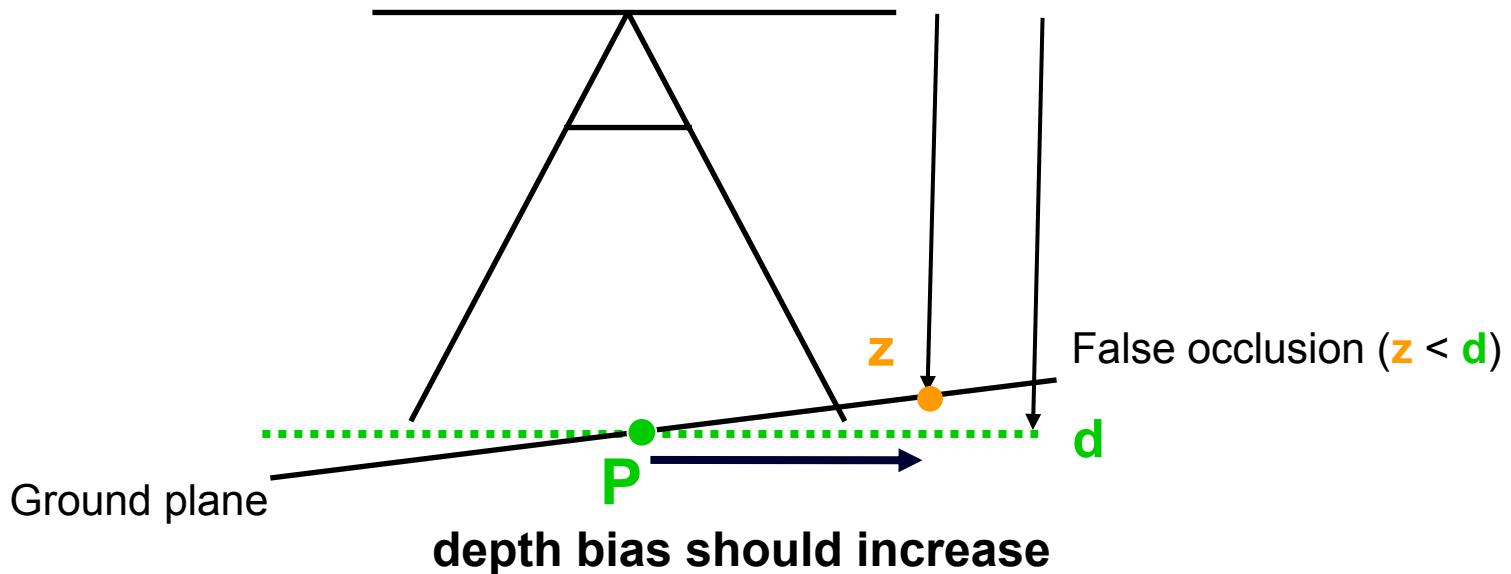


irregular sampling



PCF Self-Shadowing Issue

- Traditional depth bias: write $(z + \text{bias})$ in shadow map
 - Bias = constant bias + slope-based bias
 - Issue: huge depth biases may be required for large PCF kernels





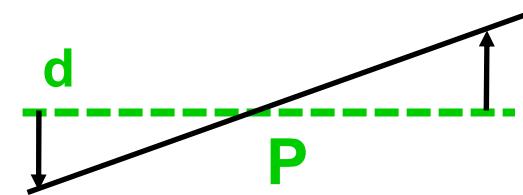
PCF Self-Shadowing Solutions

- ➊ Use depth gradient = $\text{float2}(dz/du, dz/dv)$

Make depth d follow tangent plane

$$d = d_0 + \text{dot}(\text{uv_offset}, \text{gradient})$$

[Schuler06] and [Isidoro06]



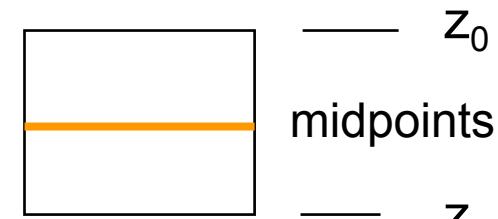
- ➋ Render midpoints into shadow map

$$\text{Midpoint } z = (z_0 + z_1) / 2$$

Requires two rasterization passes

- ➌ Depth peel two depth layers

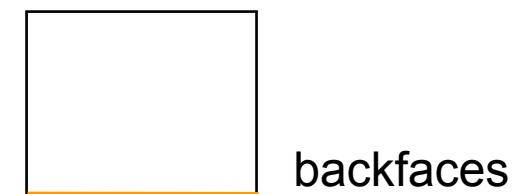
Still requires a depth bias for thin objects



- ➌ Render backfaces into shadow map

Only works for closed objects

Light bleeding for large PCF kernels

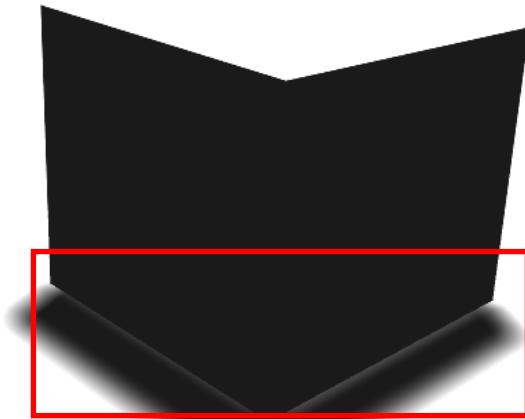




GameDevelopers
Conference 08

Backfaces and Large PCF kernels

- ⌚ Rendering backfaces into shadow map generates light bleeding for large PCF kernels
 - Not due to FP precision or shadow map resolution
 - But reverse of the surface acne issue





Non-Linearity of PCF

- ➊ PCF cannot be prefiltered as is
 - ➊ Average($d < z$) \neq (Average(z) $<$ d)
 - ➋ Filtering the depth buffer would smooth the heightfield of the shadow map
 - ➊ Does not generate soft shadows
 - ➋ May introduce artifacts
- ➋ Solutions: Approximate shadow test by a linear function which can be prefiltered
 - ➊ Goal: blurring the shadow map to generate realistic soft shadows



Outline

➊ Fixed-Size Penumbra

- ➌ PCF (Percentage Closer Filtering)
- ➌ VSM (Variance Shadow Maps)
- ➌ CSM (Convolution Shadow Maps)
- ➌ ESM (Exponential Shadow Maps)

➋ Variable-Size Penumbra

- ➌ PCSS (Percentage Closer Soft Shadows)
- ➌ PCSS + VSM/CSM
- ➌ Backprojection



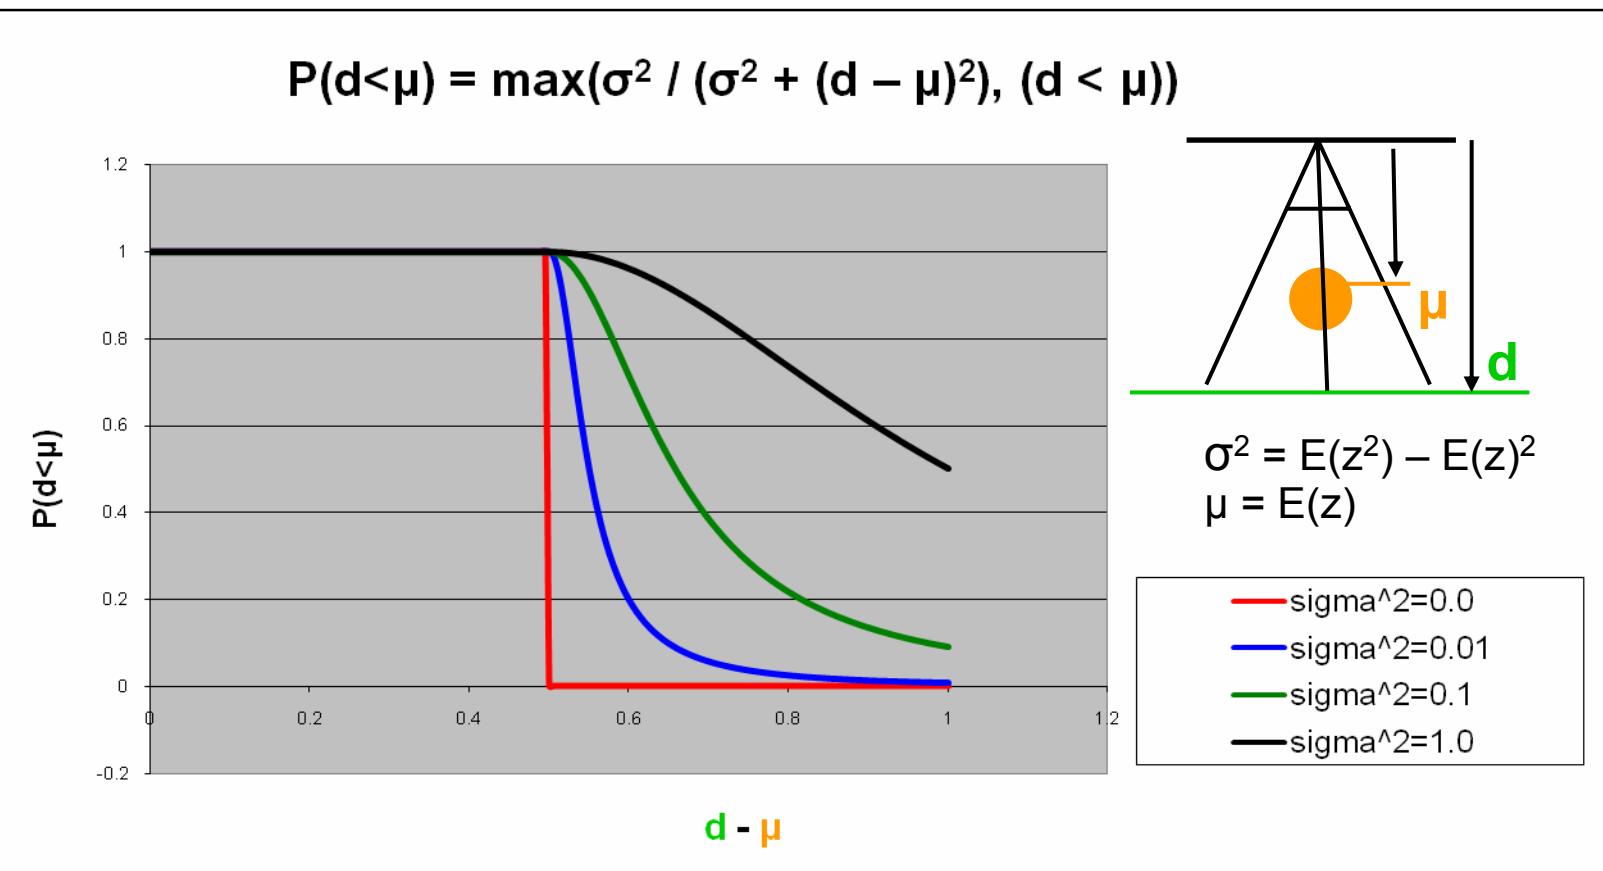
Variance Shadow Maps

- ➊ Consider depth values in the filter kernel as a depth distribution [Donnelly06] [Lauritzen07]
- ➋ Approximate the depth values in the kernel by a Gaussian distribution of mean μ and variance σ^2
$$\sigma^2 = E(z^2) - E(z)^2$$
$$\mu = E(z)$$
- ➌ Probability of being lit
 - Using Chebyshev's inequality
 - $P(d < z) \leq \max(\sigma^2 / (\sigma^2 + (d - \mu)^2), (d < \mu))$
- ➍ $E(X) = \text{average of } X \text{ in filter kernel}$
 - So the VSM (z, z^2) can be prefiltered
 - Can also be rendered using MSAA + resolve



Game Developers
Conference 08

VSM Visibility Function



- ⌚ Self-shadowing can be handled simply by clamping σ^2 to some small MinVariance parameter



VSM Rendering

- ④ Blur z and $z^2 \rightarrow E(z), E(z^2)$
- ④ Fetch $E(z), E(z^2)$ with bilinear texture lookup
 - Can also use trilinear and/or anisotropic filtering to reduce aliasing
- ④ Compute $\mu = E(z)$ and $\sigma^2 = E(z^2) - E(z)^2$
 - If ($\sigma^2 < \text{MinVariance}$) $\sigma^2 = \text{MinVariance}$
- ④ Light intensity approximation
 - $L = \max(\sigma^2 / (\sigma^2 + (d - \mu)^2), (d < \mu))$



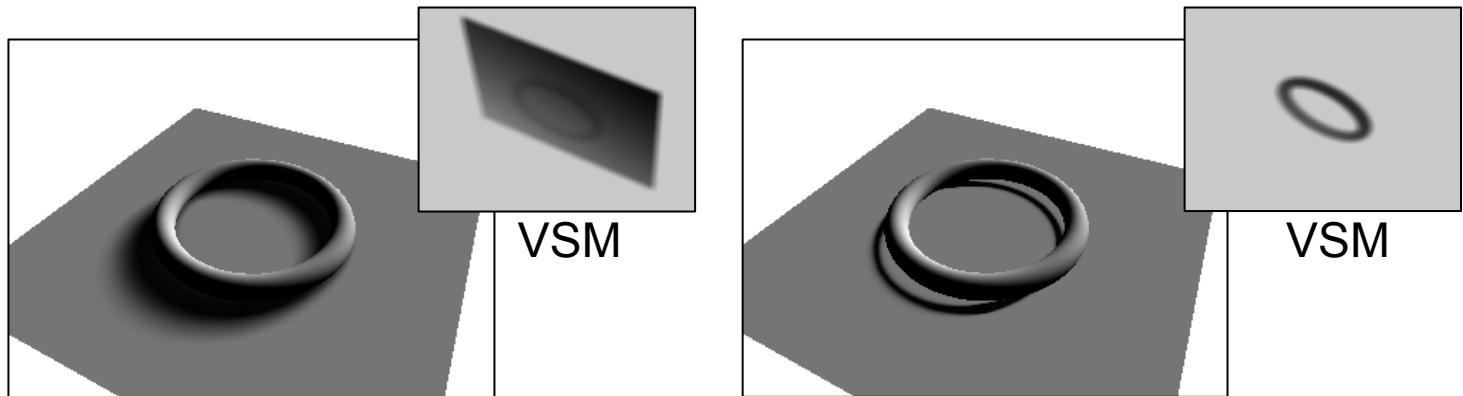
Generating the VSM

- ➊ Can use R32G32F_FLOAT format
 - Filterable format on all DX10-class graphic hardware
 - Can also use FP16 [Donnelly06]
- ➋ Render (z , z^2) to VSM texture
 - Where z is a linear distance
 - ➌ Can use an orthographic projection
 - ➌ $z = \text{mul}(\text{wPos}, \text{mObjectToLightOrtho}).z$
- ➌ Clear VSM to (MAX_Z , $\text{MAX_Z} * \text{MAX_Z}$)



Rendering everything into the VSM shadow map

- ⌚ Images with and without plane rendered in VSM



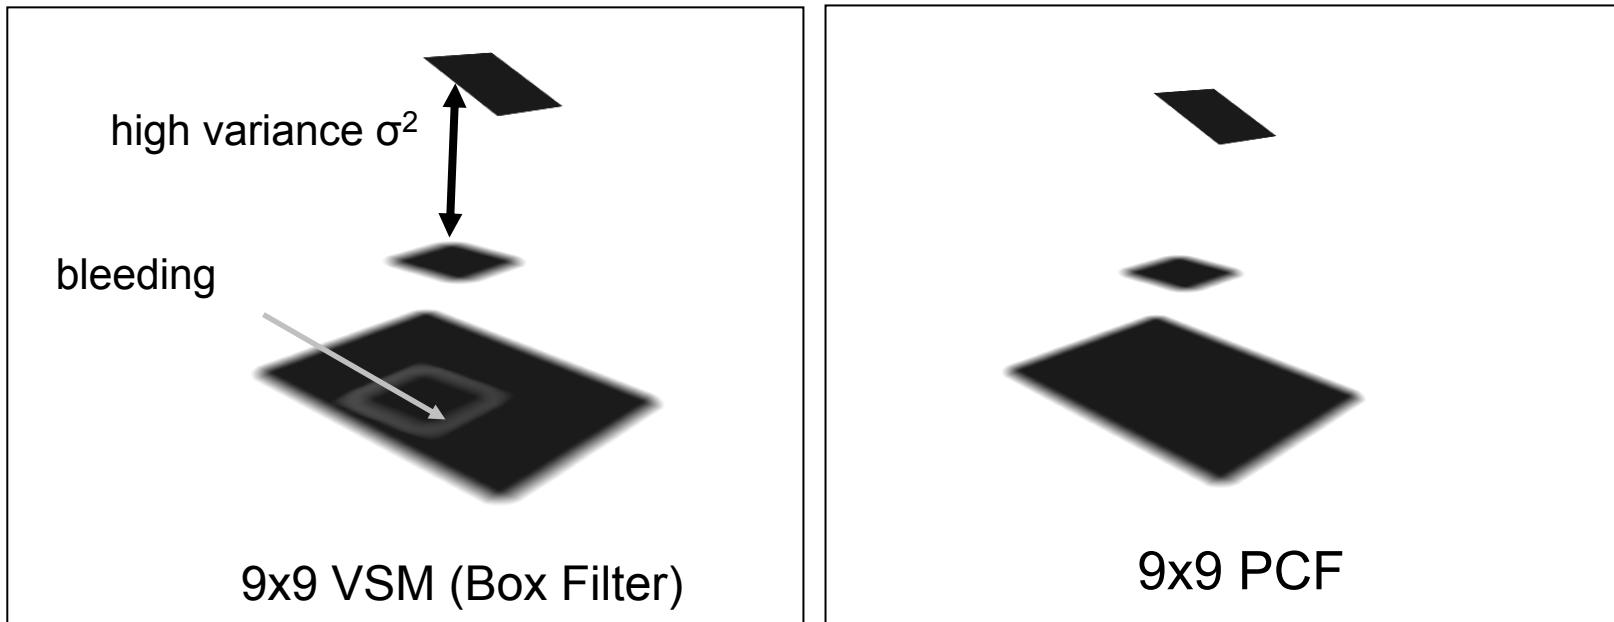
- ⌚ If receivers are not in the VSM, will blur the object texels with $z = \text{MAX_Z} \rightarrow$ missing occlusion
- ⌚ CSMs and ESMs also have this limitation
Shadows look bad when blurring shadow map without everything rendered into them



Game Developers
Conference 08

VSM Light Bleeding

Two quads floating above a ground plane



Large depth complexity in filter kernel

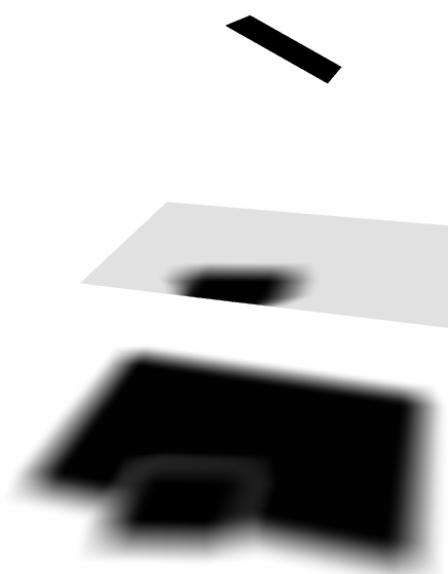
→ σ^2 is large

→ Upper bound $\sigma^2 / (\sigma^2 + (d - \mu)^2) \rightarrow 1$

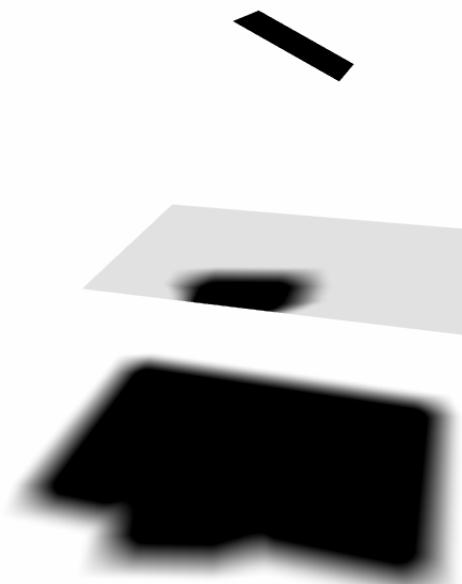


VSM Light Bleeding

- Proposed in GPU Gems 3 [Lauritzen07]
 - Use a threshold to remap shadow intensity



threshold=0



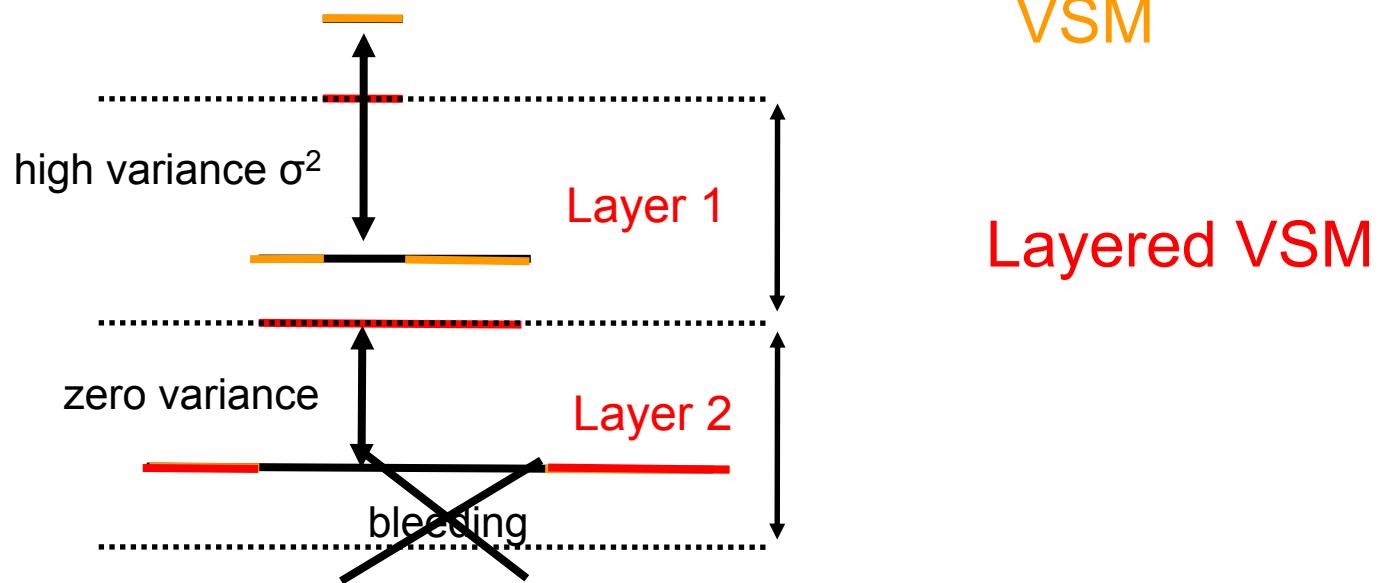
threshold=0.2



Layered VSM

⌚ [Lauritzen08] I3D 2008

Partition shadow-map frustum into multiple depth ranges
and clamp z to each range





Layered VSM

- Compared to VSMs, LVSMs
 - Reduce or eliminate light bleeding artifacts
 - Do not require 32-bit float texture filtering
 - Can still be generated in one pass (using MRTs) and need only one texture sample per pixel

Variance Shadow Map



Layered Variance Shadow Map





Outline

➊ Fixed-Size Penumbra

- ➌ PCF (Percentage Closer Filtering)
- ➌ VSM (Variance Shadow Maps)
- ➌ CSM (Convolution Shadow Maps)
- ➌ ESM (Exponential Shadow Maps)

➋ Variable-Size Penumbra

- ➌ PCSS (Percentage Closer Soft Shadows)
- ➌ PCSS + VSM/CSM
- ➌ Backprojection



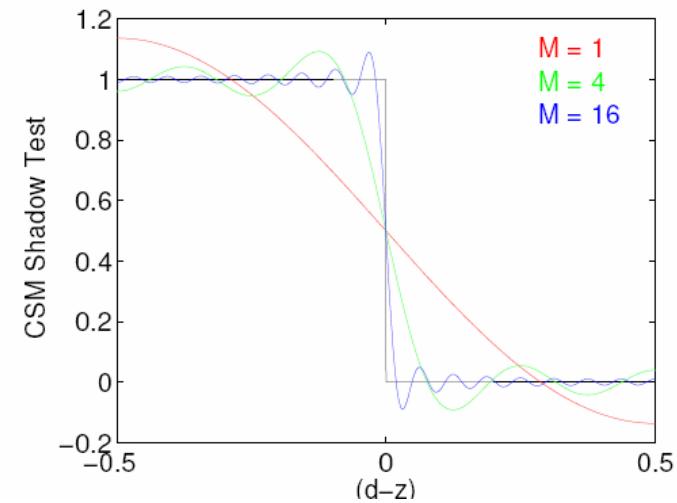
Convolution Shadow Maps

- Approximate the step function (function of depth) by 1D Fourier expansion [Annen07]

$$f(d, z) \approx \frac{1}{2} + 2 \sum_{k=1}^M \frac{1}{c_k} \cos(c_k d) \sin(c_k z) - 2 \sum_{k=1}^M \frac{1}{c_k} \sin(c_k d) \cos(c_k z)$$

Function of d
Reconstruct when
rendering shadows

Function of z
Write in shadow map
and blur it



Excerpted from
[Annen07]



Generating a CSM

- ➊ Assuming a linear depth z is available for every texel in the shadow map
 - z needs to be normalized in $[0,1]$ because the Fourier expansion is periodic
- ➋ Can generate CSM with full screen shader pass
 - Rendering to R8G8B8A8_UNORM MRTs
 - 4 MRTs ($M=8$) is usually a good tradeoff
- ➌ Background texels cleared to CSM(1.0f)



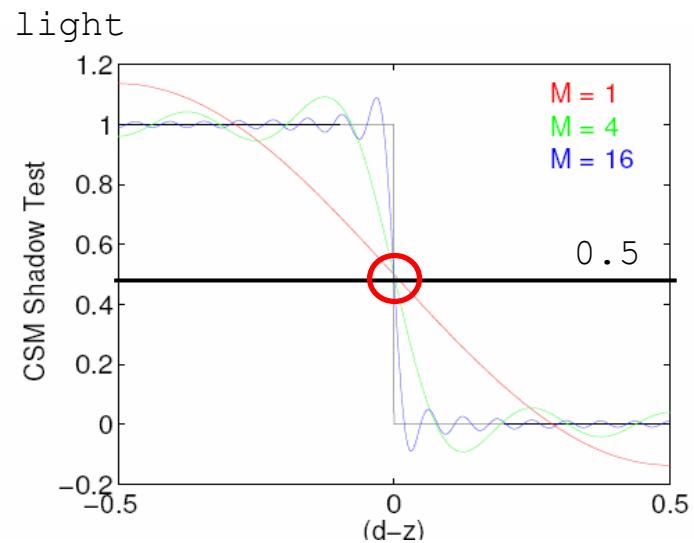
CSM Rendering

- Fetch $\cos(ck^*z)$ and $\sin(ck^*z)$ from the CSM textures and compute light intensity

$$f(d, z) \approx \frac{1}{2} + 2 \sum_{k=1}^M \frac{1}{c_k} \cos(c_k d) \sin(c_k z) - 2 \sum_{k=1}^M \frac{1}{c_k} \sin(c_k d) \cos(c_k z)$$

- Final light intensity

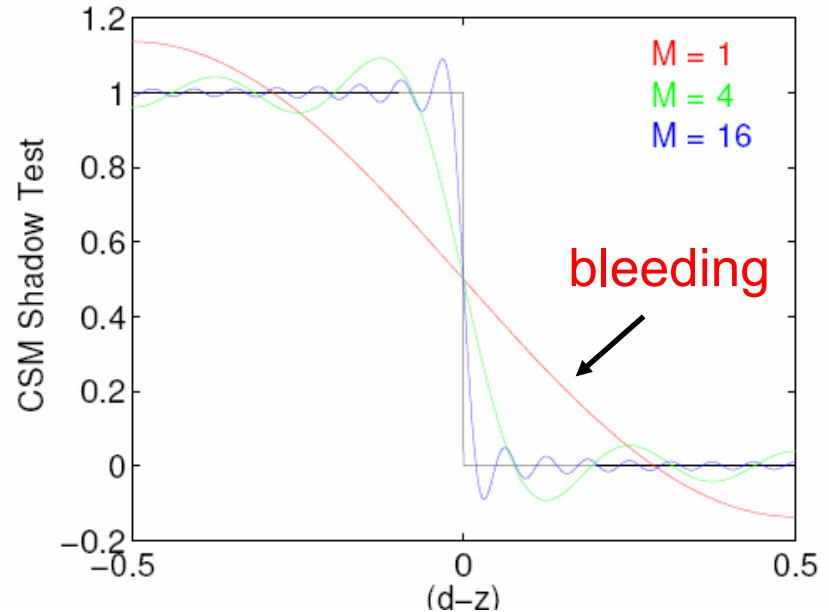
Scale by 2 otherwise $f(d,z) == 0.5$ for $(d==z)$
 $L = \text{saturate}(2.0 * f(d,z))$





CSM Light Bleeding

- ➊ Limited bleeding may be seen as a feature (deep shadow look)
- ➋ Light bleeding when should be in shadow ($d>z$) but light > 0

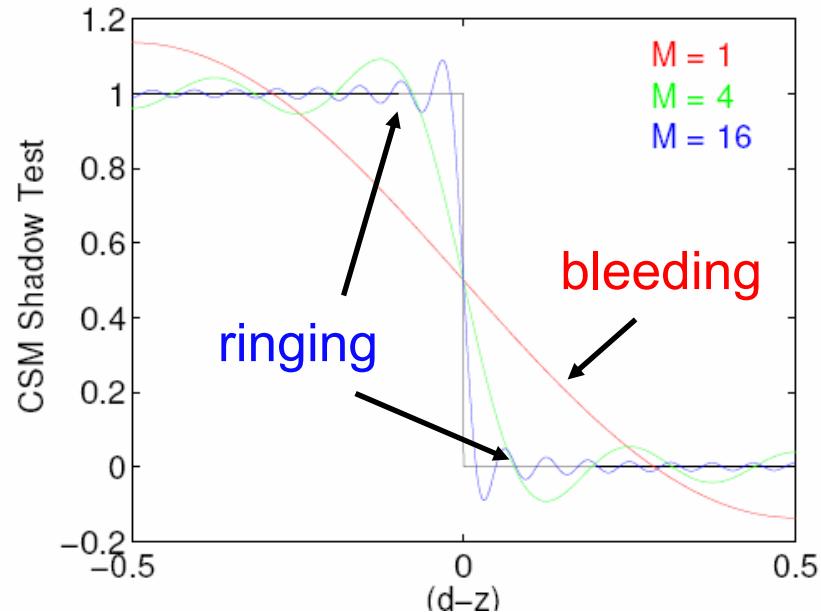




CSM Ringing

- ➊ Gibbs phenomenon
High order Fourier terms generate ringing
- ➋ Ringing workaround
Attenuate terms based on their order (k)

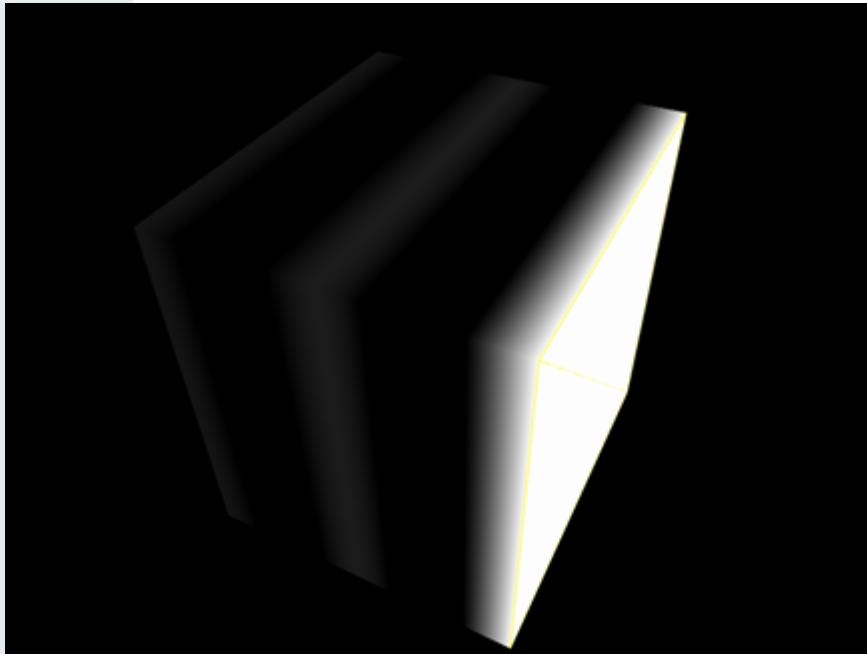
$$f(d, z) \approx \frac{1}{2} + 2 \sum_{k=1}^M \left[\frac{1}{c_k} \cos(c_k d) \sin(c_k z) - 2 \sum_{k=1}^M \frac{1}{c_k} \sin(c_k d) \cos(c_k z) \right] * \exp(-\alpha(k/M)^2)$$



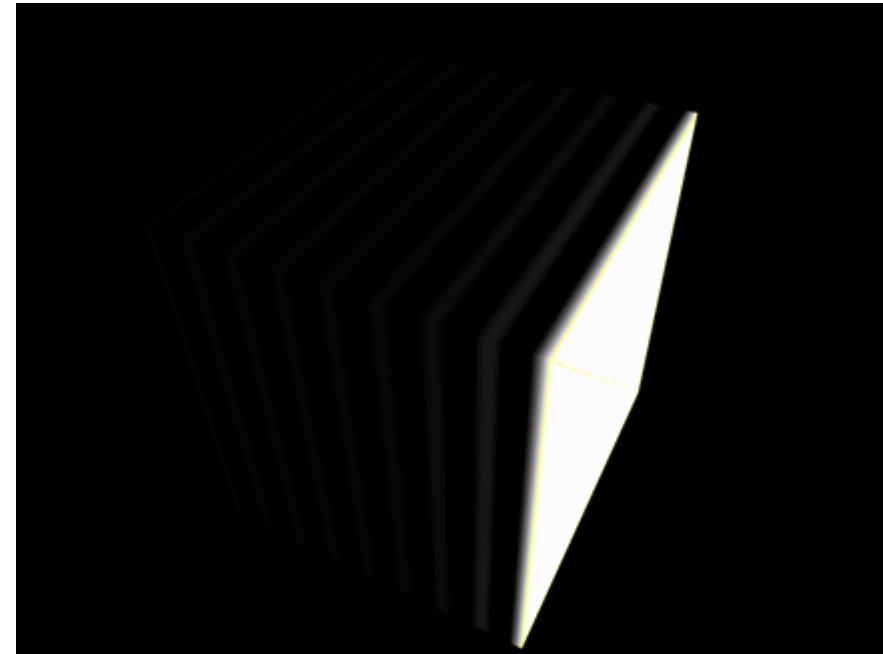


CSM Ringing

- Light source facing inwards the cube



2 RGBA8

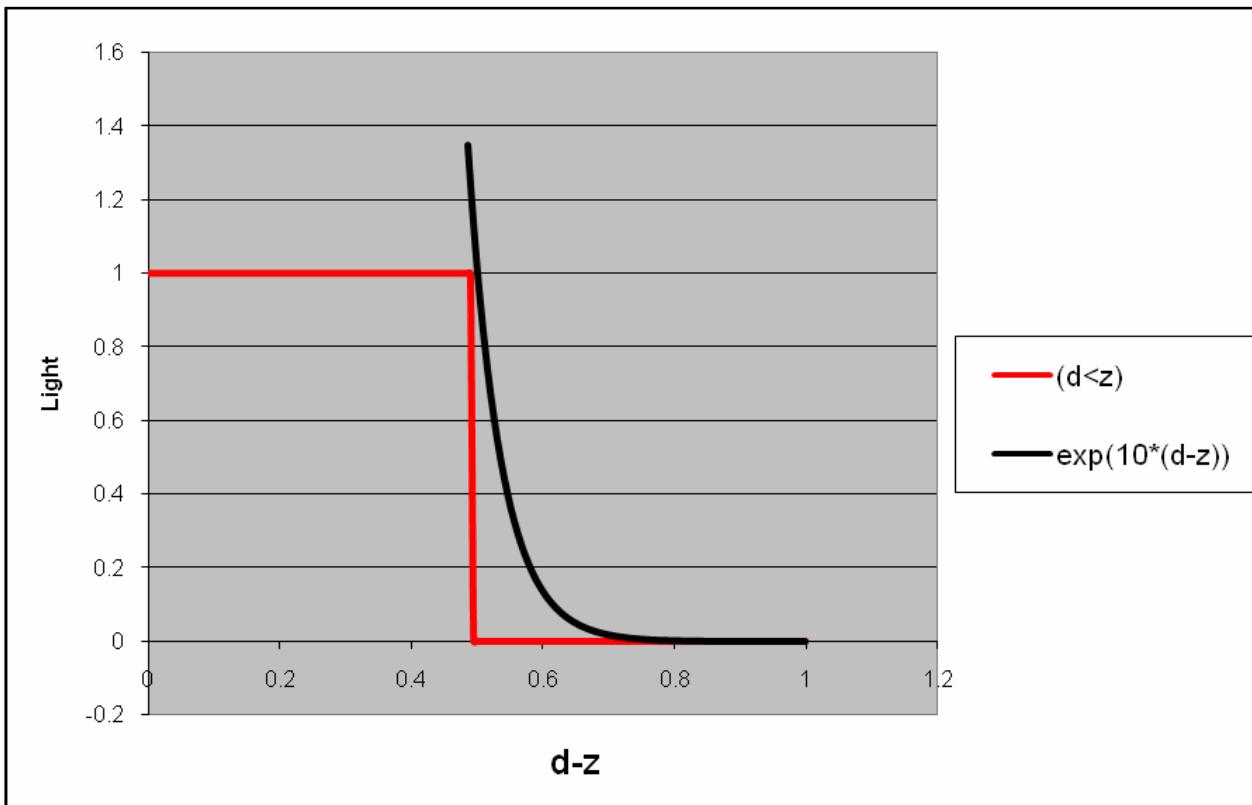


8 RGBA8



Exponential Shadow Maps

- ⌚ [Salvi08] ShaderX⁶
- ⌚ Approximate step function ($z-d > 0$) by
$$\exp(k^*(z-d)) = \exp(k^*z) * \exp(-k^*d)$$





ESM Parameter Tuning

$$L = \exp(k*z) * \exp(-k*d)$$



$k=10$



$k=30$



Deep Shadow Look

- ⦿ Shadows without any shading
Soft shadow test gives translucent look



CSM



VSM



ESM



VSM/CSM/ESM

Parameters

VSM: MinVariance, BleedingReductionFactor

CSM: NumTextures, AbsorptionFactor

ESM: ScaleFactor

Storage

VSM	CSM	ESM
R32G32	N * (R8G8B8A8)	R32

For CSM, N>=4 is usually required to avoid bleeding

The less storage, the faster the prefiltering



VSM/CSM/ESM

GameDevelopers
Conference 08



VSM – 87 fps



CSM – 38 fps



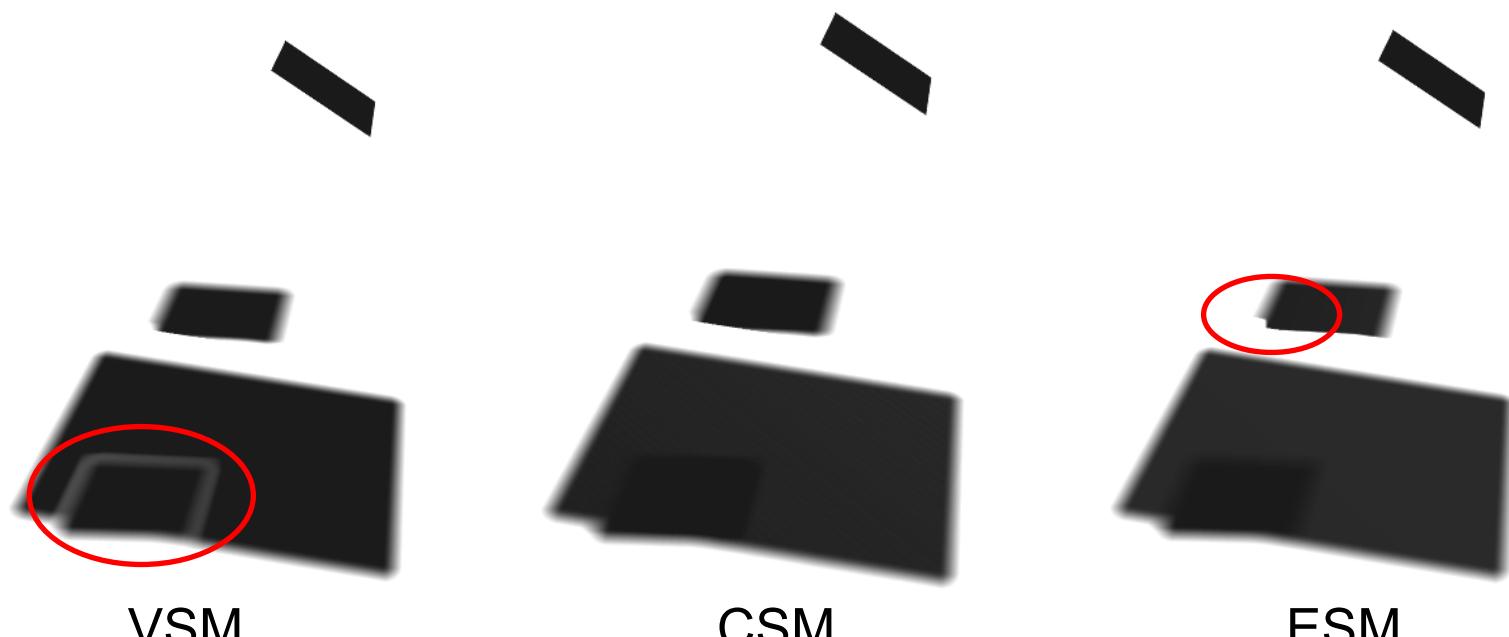
ESM – 93 fps

8800 GTX Ultra



VSM/CSM/ESM

- ⌚ 2 quads floating above a ground plane



- ⌚ The minor ESM artifacts can be handled with another parameter which overdarkens the shadows [Salvi08]



Outline

➊ Fixed-Size Penumbra

- ➌ PCF (Percentage Closer Filtering)
- ➌ VSM (Variance Shadow Maps)
- ➌ CSM (Convolution Shadow Maps)
- ➌ ESM (Exponential Shadow Maps)

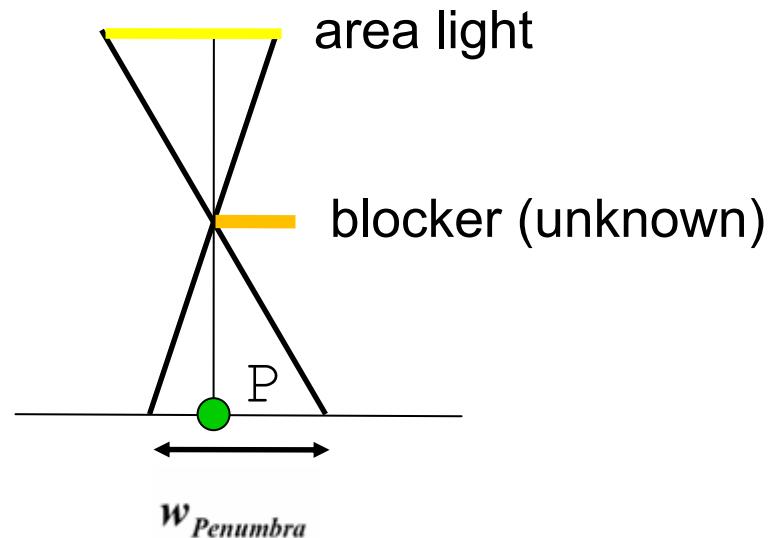
➋ Variable-Size Penumbra

- ➌ PCSS (Percentage Closer Soft Shadows)
- ➌ PCSS + VSM/CSM
- ➌ Backprojection



Percentage Closer Soft Shadows

- ⦿ PCSS [Fernando05]
- ⦿ Assume a square light centered at the shadow map center
- ⦿ Assuming some parallel blocker to receiver
 - Compute penumbra width using similar triangles



$$w_{Penumbra} = \frac{(d_{Receiver} - d_{Blocker}) \cdot w_{Light}}{d_{Blocker}}$$



PCSS Overview

Step 1: Blocker Search

- ➊ Sample the depth buffer using point sampling
- ➋ Average all blockers with $(\text{depth} + \text{bias} < \text{receiver})$ in search region / kernel
- ➌ Early out if no blocker found

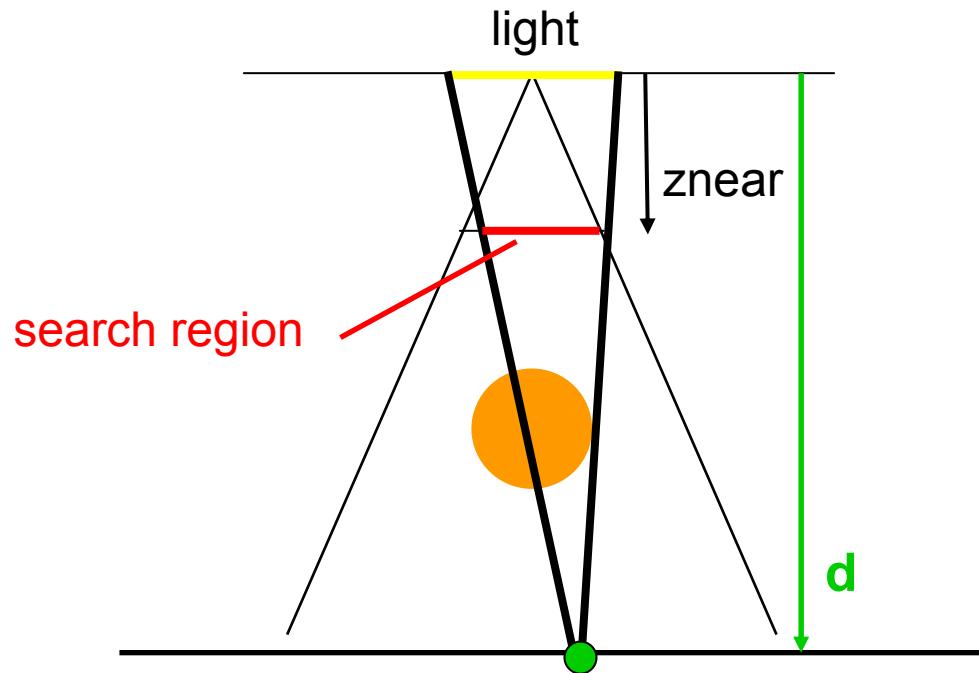
Step 2: Filtering

- ➊ Use filter radius from step 1
- ➋ Clamp filter width to be $\geq \text{MinRadius}$ for antialiasing
- ➌ Filter the shadow map with PCF or VSM/CSM/ESM



Search Region

- Where to find blockers?
- Conservative search radius using similar triangles

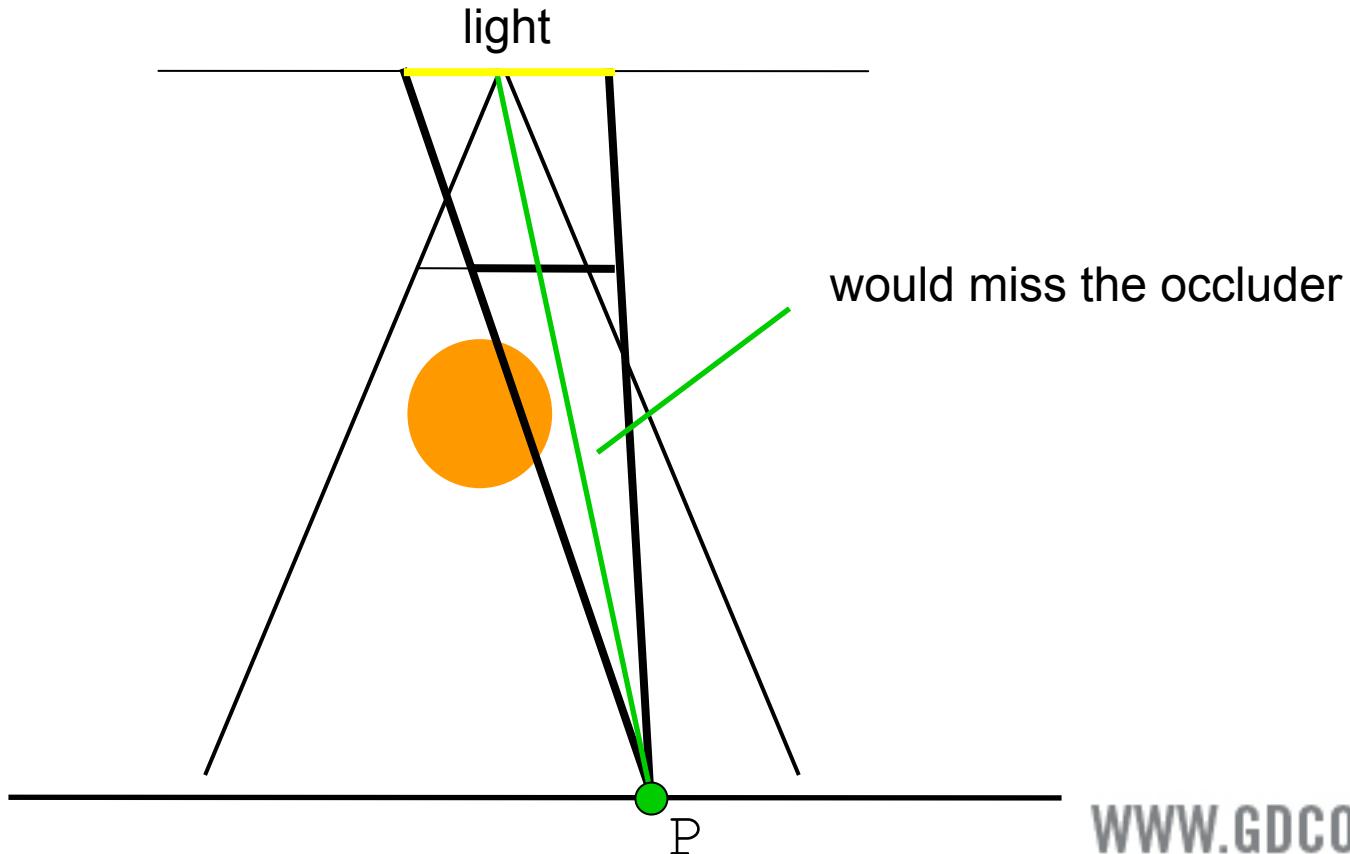


$$\text{LightRadius} / d = \text{SearchRadius} / (d - znear)$$



Blocker Search

- Why not doing just one sample?





Blocker Search

- ⌚ The more samples in the blocker search, the less noisy artifacts in the soft shadows
In practice, 4x4 or 5x5 samples is sufficient



Blocker Search with 3x3 taps



Blocker Search with 5x5 taps



Hellgate_DX10_vs
DX9_Large.jpg...

PCSS In Hellgate: London



PCSS



PCF

16 POINT taps for the blocker search
16 PCF taps for the PCF filtering



Outline

➊ Fixed-Size Penumbra

- ➌ PCF (Percentage Closer Filtering)
- ➌ VSM (Variance Shadow Maps)
- ➌ CSM (Convolution Shadow Maps)
- ➌ ESM (Exponential Shadow Maps)

➋ Variable-Size Penumbra

- ➌ PCSS (Percentage Closer Soft Shadows)
- ➌ PCSS + VSM/CSM
- ➌ Backprojection



PCSS + VSM/CSM/ESM

- ⊕ Compute filter size using blocker search
- ⊕ Solutions to prefilter the shadow map to support fast blurs of variable size
 - ⊕ Summed Area Tables (SAT)
 - ⊕ Fast Box Filters
- ⊕ Mipmapping
 - ⊕ $\text{Level}[i] = \text{Blur}(\text{Level}[i-1])$



Summed Area Tables

- ➊ For a 1D texture, each texel $S[n]$ of the SAT is
$$S[n] = \text{Sum}(T[i], i=0..n)$$
- ➋ Average in interval $(i,j]$ is
$$(S[j]-S[i]) / (j-i)$$
- ➌ Also works in 2D
With more 4 fetches instead of 2
- ➍ Fast algorithm for building SATs on GPUs [Hensley05]
Recursive doubling
For $N \times N$ shadow map, $\log(N)$ passes



Summed Area Shadow Maps

⌚ Using UINT32

- ⌚ Advantages: more bits available than FP32, and support arbitrary shadow map resolution
- ⌚ Drawback: need to implement bilinear filtering in shader
- ⌚ See GPU Gems 3 chapter for details [Lauritzen07]

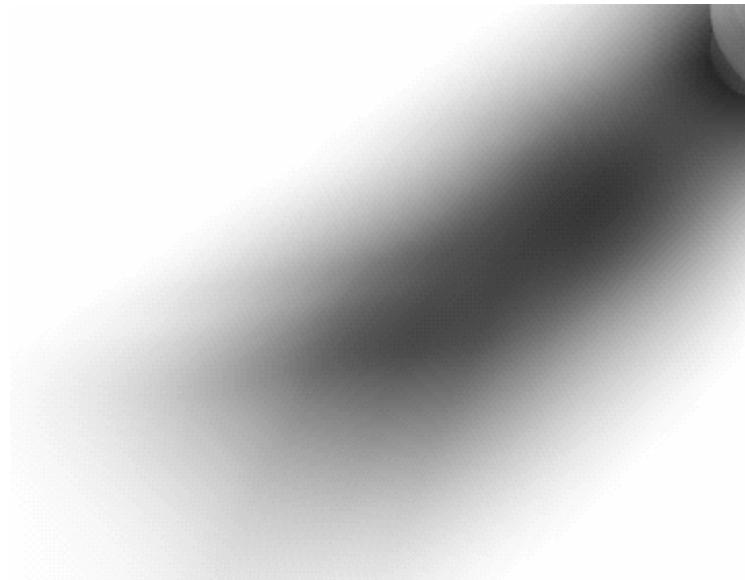
⌚ Best fitting formats for DX10 GPUs

- ⌚ For SAVSM - R32G32_UINT
- ⌚ For SACSM - multiple R32G32B32A32_FLOATs
- ⌚ For ESM (storing $\exp(k^*z)$) - R32_UINT



SAVSM vs MIPCSM

- ➊ Disadvantage of SAT: can only do box filter
- ➋ Mipmap with recursive blurs looks more like a Gaussian



SACSM



MIPCSM



SATs with variable size kernel

- ➊ Rounding the filter radius generates banding
`float2 moments = FilterMomentsBilinear(uv, round(filterRadius));`

- ➋ Better quality: trilinear filtering
`float2 moments0 = FilterMomentsBilinear(uv, floor(filterRadius));`
`float2 moments1 = FilterMomentsBilinear(uv, ceil(filterRadius));`
`float2 moments = lerp(moments0, moments1, frac(filterRadius));`



SATs for VSM/CSM/ESM

- Ⓐ Cost of building a SAT (on G80 Ultra)
 - SAVSM
 - Ⓐ 0.9 ms for 512^2 , 3.3 ms for 1024^2
 - SACSM with 16 FLOATs per texel
 - Ⓐ 6.3 ms for 512^2 and 27.5 ms for 1024^2
 - SAESM
 - Ⓐ 0.6 ms for 512^2 , 2.2 ms for 1024^2
- Ⓐ CSM is not well suited for SATs
 - Mipmaps are a better fit for them



MIPCSM

- ➊ Build a mipmap pyramid on top of CSM
 - Implemented as texture array with mip levels
 - How to generate the mipchain
 - ➋ Default 2x2 mipchain generates aliasing
 - ➋ Use a larger box filter (at least 7x7) to generate each LOD, double spacing every level
 - ➋ Lookup CSM using trilinear filtering
- ➋ Advantages of MIPCSM over SACSM
 - ➋ Can use RGBA8 instead of RGBA32F
 - ➋ Much faster to build than SACSM



MIPCSM

Game Developers
Conference 08

- To reduce issues to discretized blurs, do a small blur pre-pass over the CSM



without pre-blur



with 5x5 pre-blur



Outline

➊ Fixed-Size Penumbra

- ➌ PCF (Percentage Closer Filtering)
- ➌ VSM (Variance Shadow Maps)
- ➌ CSM (Convolution Shadow Maps)
- ➌ ESM (Exponential Shadow Maps)

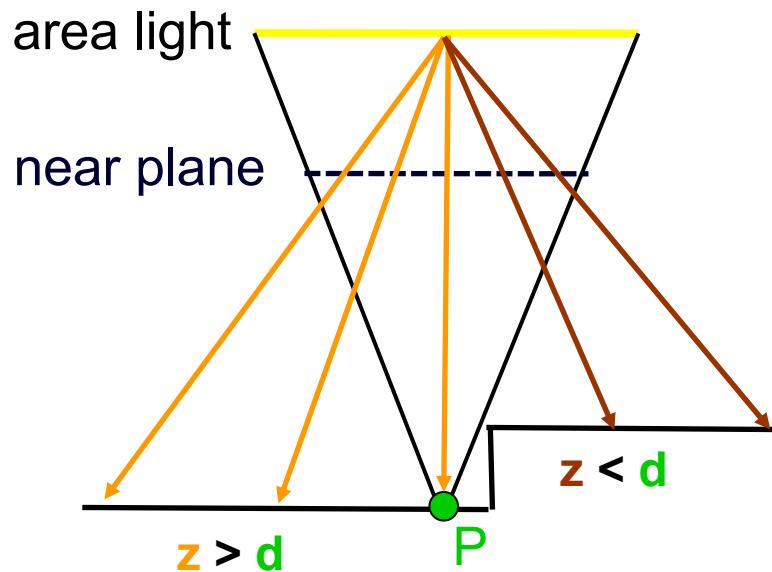
➋ Variable-Size Penumbra

- ➌ PCSS (Percentage Closer Soft Shadows)
- ➌ PCSS + VSM/CSM
- ➌ Backprojection

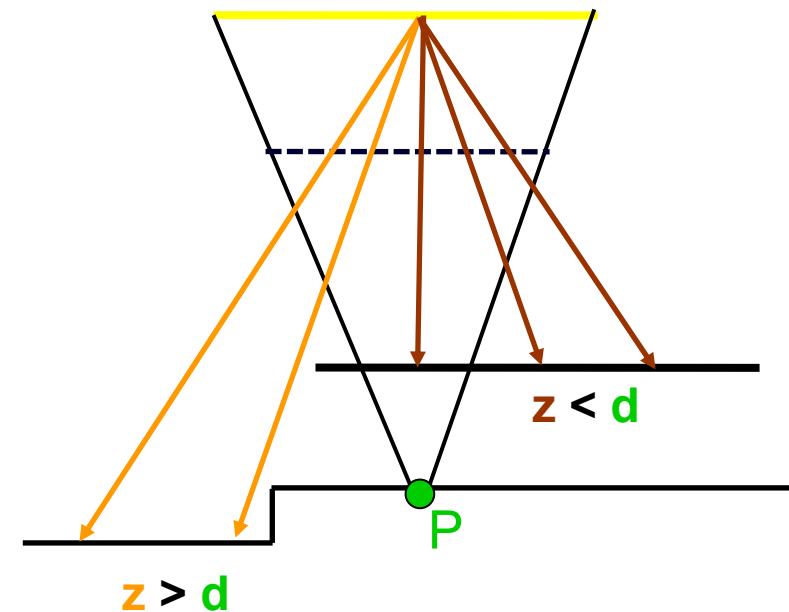


GameDevelopers
Conference 08

Why PCF is wrong for soft shadows



PCF: $\text{Shadow}(P) > 0$
Ground Truth: $\text{Shadow}(P) == 0$

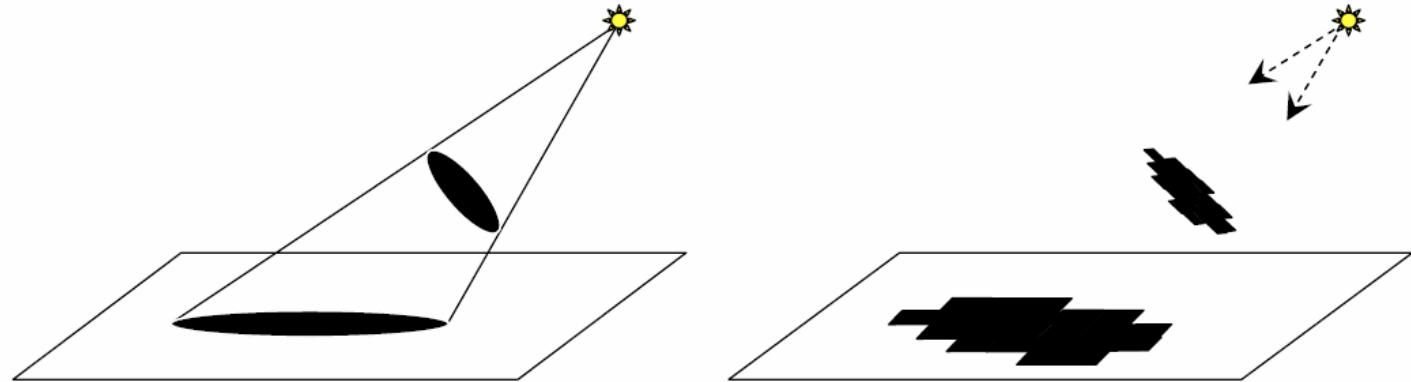


PCF: $\text{Shadow}(P) < 1$
Ground Truth: $\text{Shadow}(P) == 1$



Unprojecting the shadow map

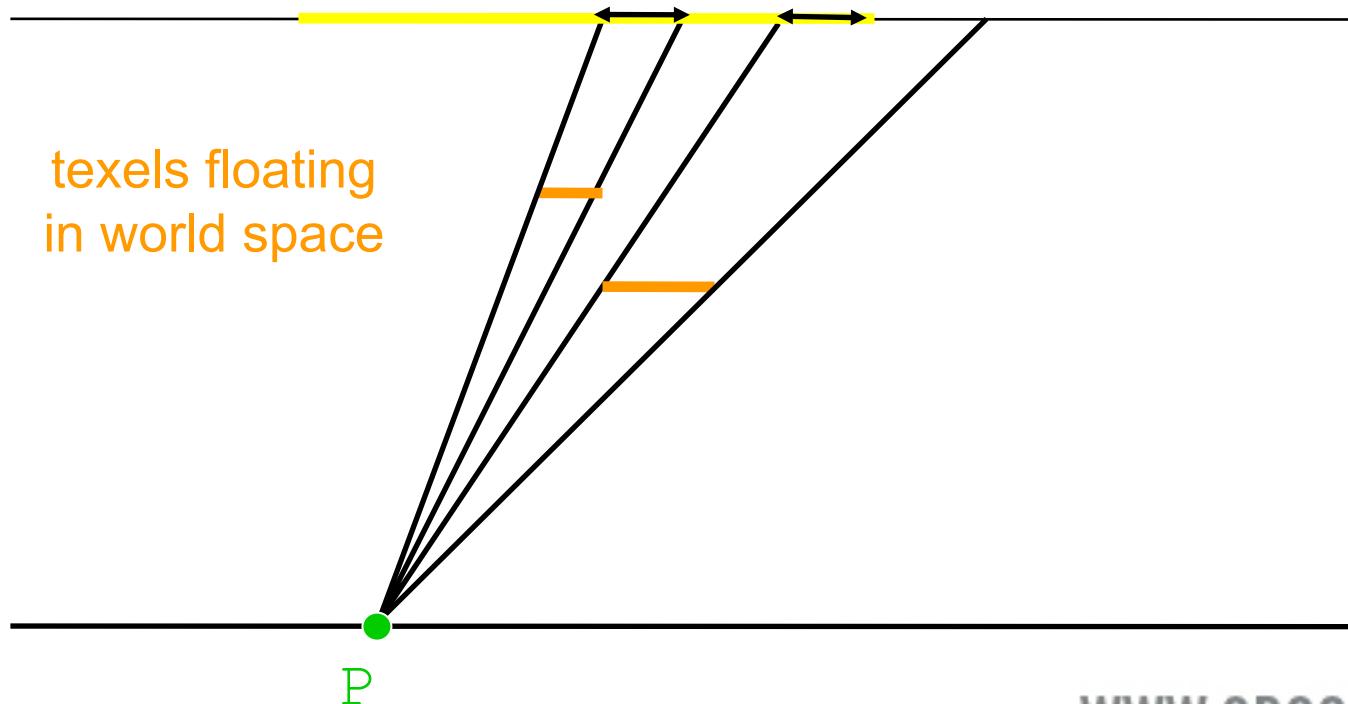
- ⌚ Unproject shadow map texels into world space [Atty05] [Guennebaud06]
- ⌚ Discretized approximation of the scene geometry





Backprojecting texels

- ➊ Backproject texels onto light source
- ➋ Clamp to light borders and accumulate area





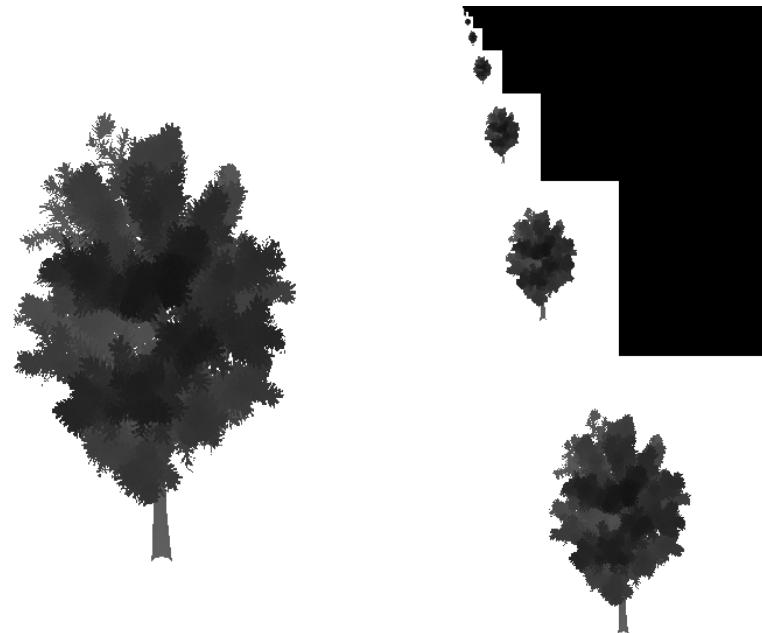
Min-Max Mipmap Shadow Map

- ⌚ Can represent shadow-map data in hierarchical fashion [Dmitriev07]
 - Hierarchical traversal allows for efficient pruning of subtrees
- ⌚ Build a hierarchical shadow map first
 - ⌚ R32G32_FLOAT mipmap
 - ⌚ Each level stores min and max z of 2x2 subpixels
- ⌚ Links
 - [SDK sample](#)
 - [GDC 2007 talk](#)



Flattening the mipmap

- ➊ Generate mipmap and flattened 2D texture
 - ➋ Load() from texture without mipmaps is faster because it guarantees that all pixels access the same mipmap level.

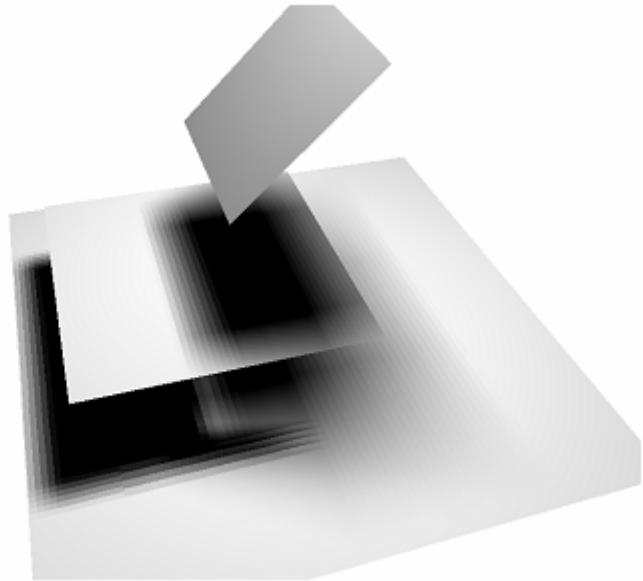
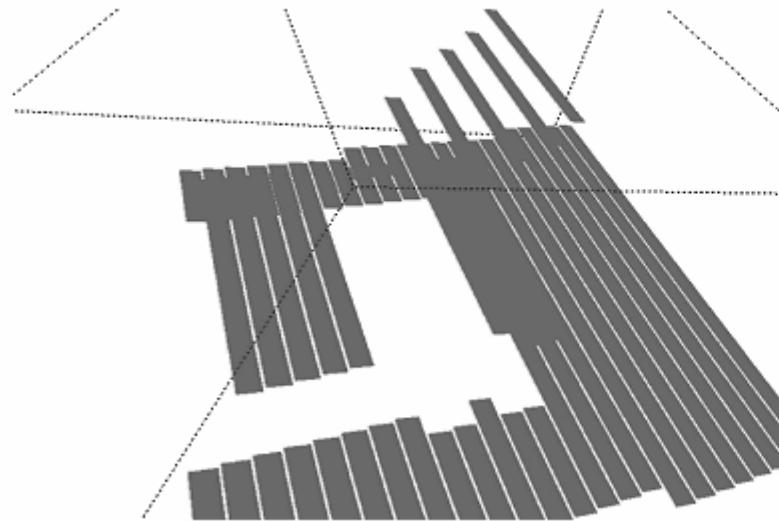


Flattened mipmap



Light Bleeding

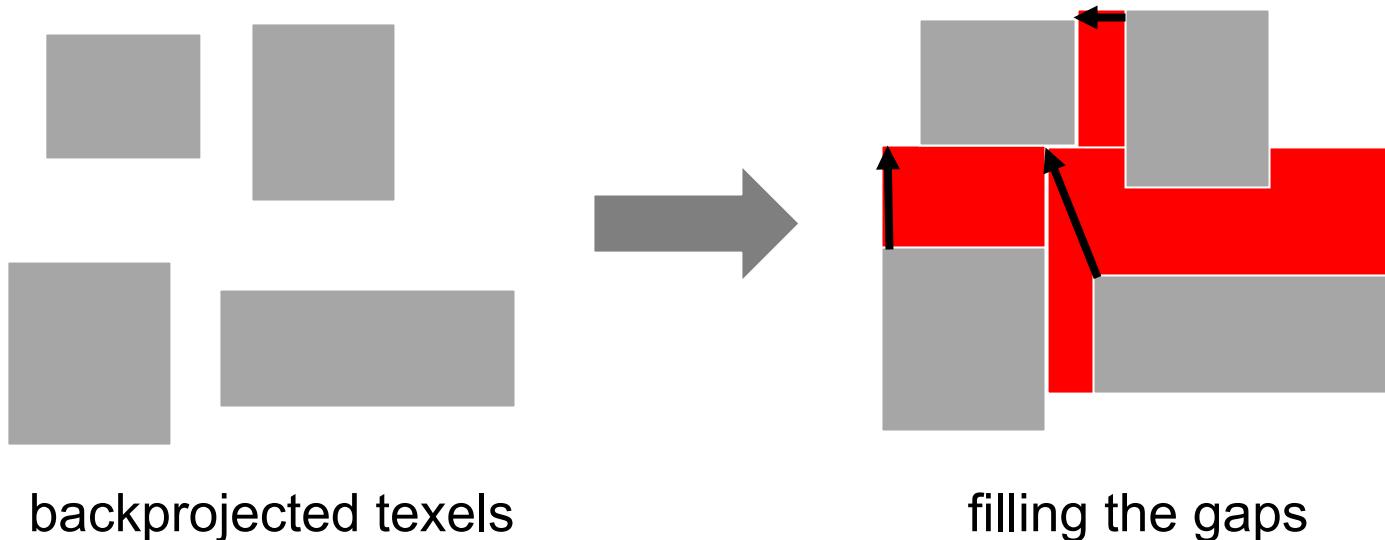
- Light bleeding is caused by gaps between micropatches





Gap Filling

- Extend backprojected quads on the sides to touch their left and top neighbors and the diagonal
Similar to [Guennebaud06]



PCSS with PCF

Frame Rate: 120 fps (8.3 ms)
Shadow Map Generation: 1.4 ms

Blocker Search: 5x5 taps
PCF: 5x5 taps

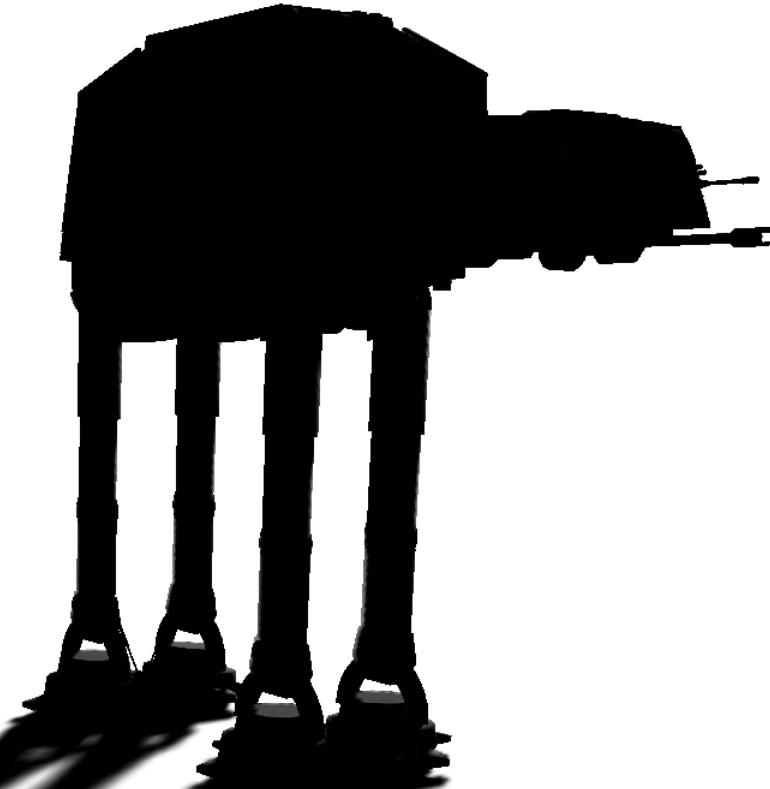


Image: 1600x1200
Shadow Map: 1024²
Triangle Count: 211k
GeForce 8800 GT

PCSS with PCF

Frame Rate: 96 fps (10.4 ms)
Shadow Map Generation: 1.4 ms

Blocker Search: 5x5 taps
PCF: 9x9 taps

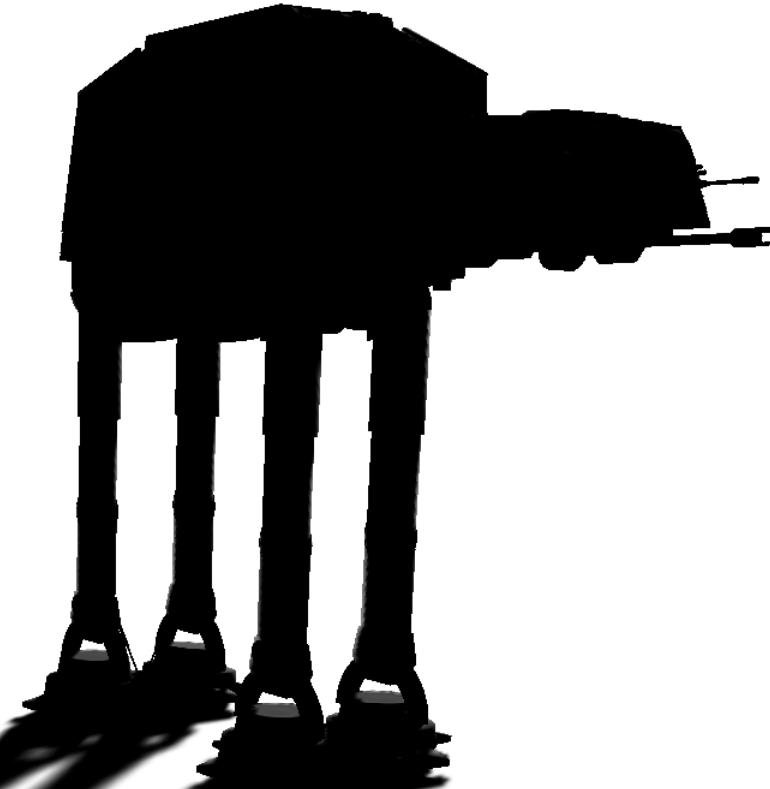


Image: 1600x1200
Shadow Map: 1024^2
Triangle Count: 211k
GeForce 8800 GT

PCSS with PCF

Frame Rate: 54 fps (18.5 ms)
Shadow Map Generation: 1.4 ms

Blocker Search: 5x5 taps
PCF: 17x17 taps

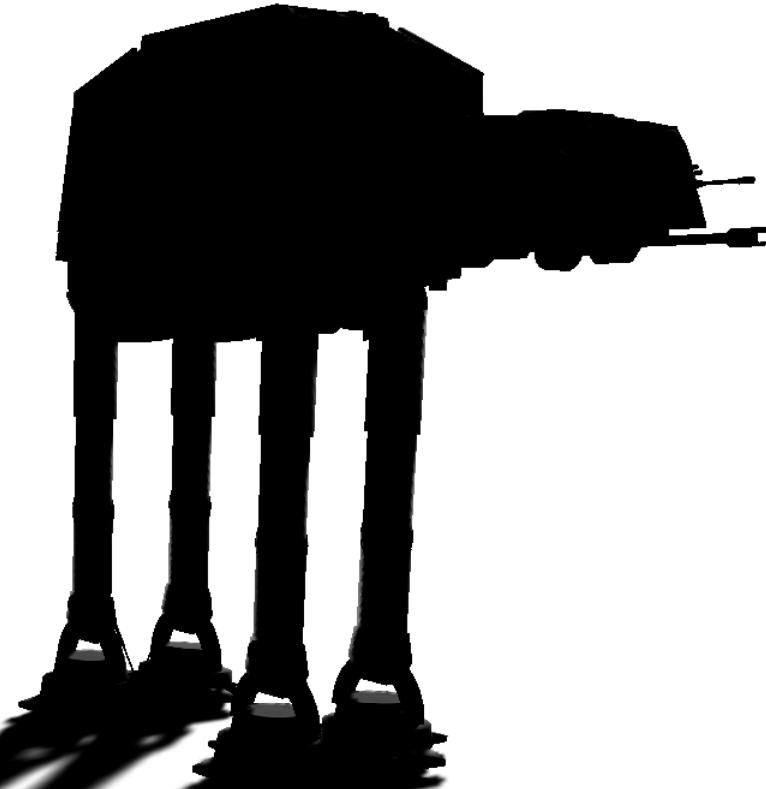


Image: 1600x1200
Shadow Map: 1024²
Triangle Count: 211k
GeForce 8800 GT

PCSS + SAVSM

Frame Rate: 68 fps (14.7 ms)
SAVSM: 6.5 ms total (44%)

Blocker Search: 5x5 taps
UINT32 SAVSM

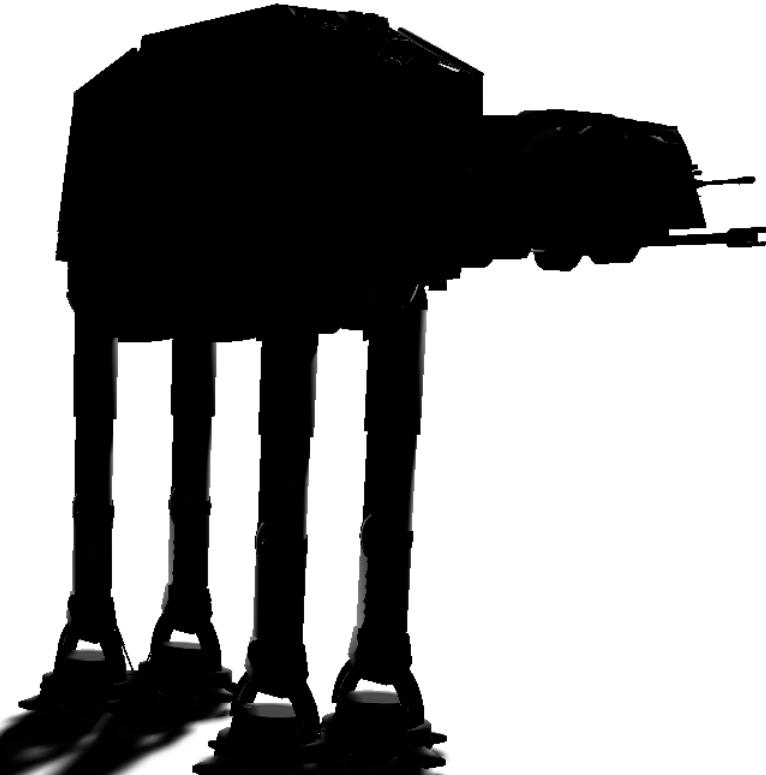


Image: 1600x1200
Shadow Map: 1024²
Triangle Count: 211k
GeForce 8800 GT

PCSS + MIPCSM

Frame Rate: 57 fps (17.5 ms)

MIPCSM: 9.4 ms total (54%)

Blocker Search: 5x5 taps

MIPCSM: 4 RGBA8

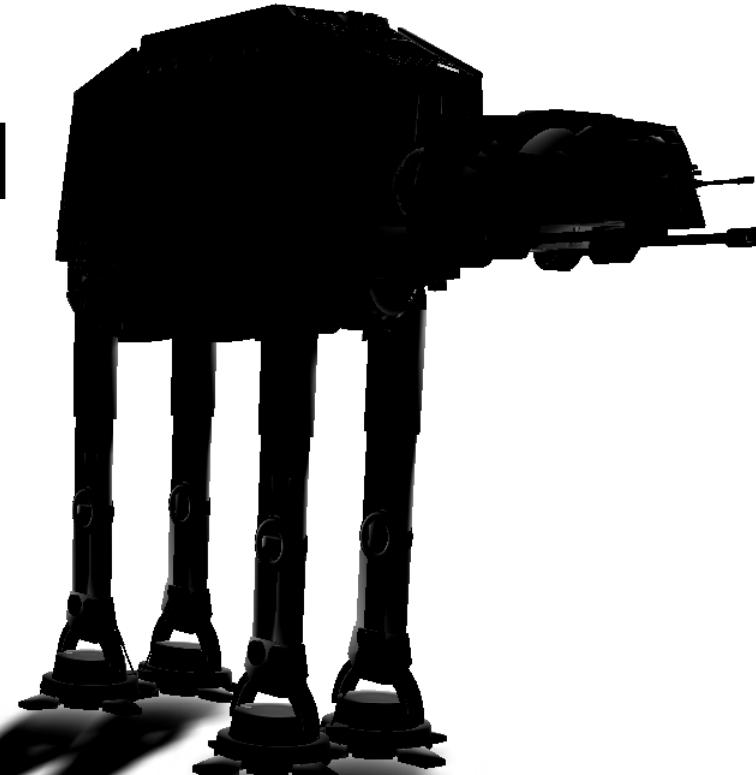


Image: 1600x1200
Shadow Map: 1024²
Triangle Count: 211k
GeForce 8800 GT

Hierarchical Backprojection

Frame Rate: 11 fps (91 ms)

Min-Max Mipmap: 2.6 ms (3%)

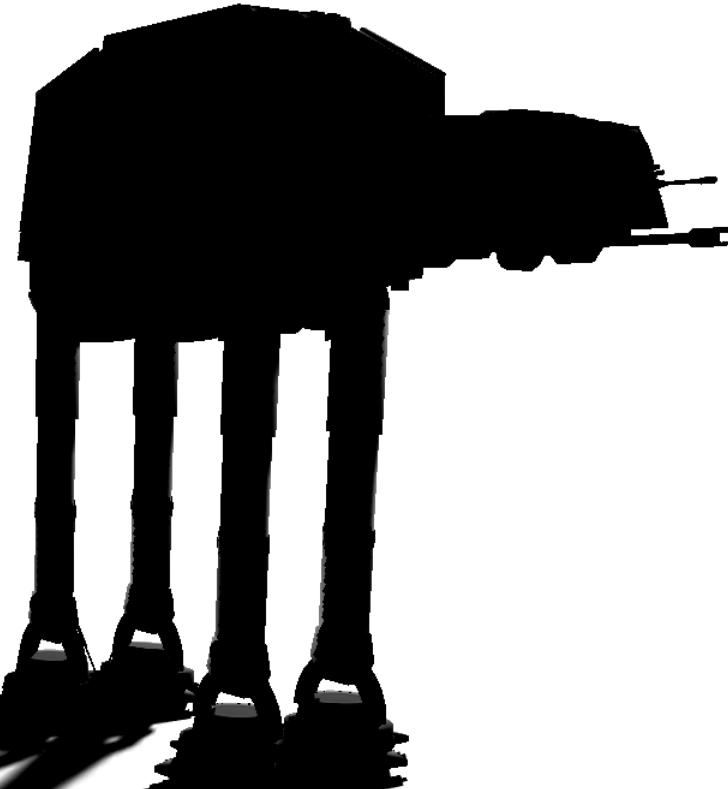
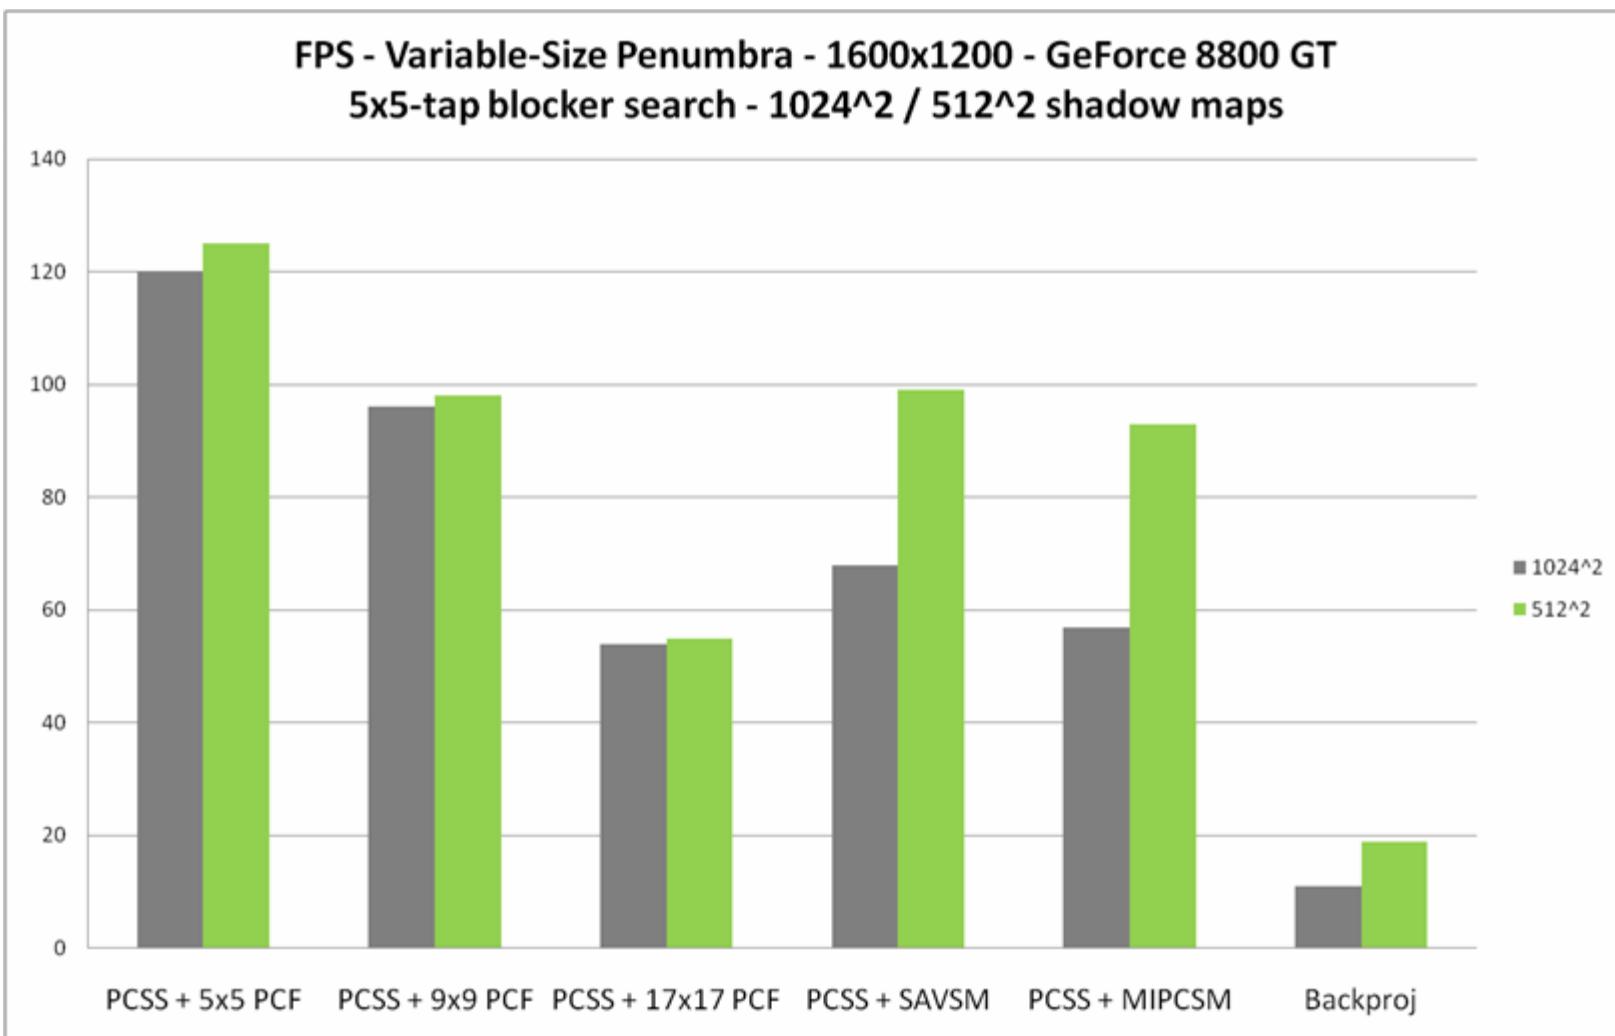


Image: 1600x1200
Shadow Map: 1024²
Triangle Count: 211k
GeForce 8800 GT



Performance Comparison





Conclusions

- ⌚ VSM/CSM/ESM all have some sort of light bleeding
 - ⌚ There are ways to control it
 - ⌚ VSM: threshold and/or layered VSM [Lauritzen08]
 - ⌚ CSM: more coefficients [Annen07]
 - ⌚ ESM: scale factor [Salvi08]
- ⌚ Percentage Closer Soft Shadows (PCSS) is good stuff
 - ⌚ Used in shipping game (Hellgate: London)
 - ⌚ See our latest [PCSS whitepaper](#) (February 2008)
 - ⌚ Two step process:
 - ⌚ 1. compute filter size using point sampling
 - ⌚ 2. filter the shadow map using PCF



Conclusions

- ⌚ Can swap PCF filtering with other filtering methods
 - PCSS + SAVSM works great
 - ⌚ See GPU Gems 3 Chapter for SAVSM (including source code)
 - ⌚ Can reduce light bleeding using layered VSMs
- ⌚ Ideas for accelerating the blocker search
 - Min-max mipmap hierarchical traversal
 - Summed Area Table in the case where the shadow receivers are not in the PCSS shadow map
- ⌚ Shadow map generation recommendation
 - Use a separate stream for the depth pre-pass with the minimum set of vertex attributes



Acknowledgments

⦿ NVIDIA

Kevin Myers
Randy Fernando
Kirill Dmitriev
Yury Uralsky
Per Vognsen

- ⦿ Marco Salvi (ESM)
- ⦿ Andrew Lauritzen (VSM)



Resources

Implementation details

PCSS

- ⌚ [PCSS integration guide](#) (February 2008)
- ⌚ DirectX 10 sample to be released soon

Hierarchical backprojection

- ⌚ [SoftShadows](#) sample in our SDK10

SAVSM

- ⌚ GPU Gems 3 chapter

Models

AT-AT: Brad Blackburn / www.scifi3d.com

Trees: Generated using [Dryad](#)

Buddha: Stanford Mesh Repository



References

- ④ [Salvi08] Marco Salvi, “Rendering Filtered Shadows with Exponential Shadow Maps”, ShaderX6
- ④ [Lauritzen08] Andrew Lauritzen, Michael McCool, “Layered Variance Shadow Maps”, I3D 2008 (poster)
- ④ [Annen07] T. Annen, T. Mertens, P. Bekaert, H.-P. Seidel, J. Kautz, “Convolution Shadow Maps”, EGSR, June 2007
- ④ [Dmitriev07] Kirill Dmitriev, Yury Uralsky “Soft shadows using hierarchical min-max shadow maps”, GDC 2007
- ④ [Schuler06] Christian Schueler, “Normal Mapping without Pre-Computed Tangents”, ShaderX5
- ④ [Isidoro06] John Isidoro, “Shadow Mapping: GPU-based Tips and Techniques”, GDC 2006



References

- ⌚ [Donnelly06] William Donnelly, Andrew Lauritzen, “Variance Shadow Maps, I3D 2006
- ⌚ [Guennebaud06] G. Guennebaud, L. Barthe, M. Paulin, “Real-time Soft Shadow Mapping by Backprojection”, EGSR 2006
- ⌚ [Fernando05] Randy Fernando, “Percentage-Closer Soft Shadows”, SIGGRAPH 2005
- ⌚ [Hensley05] J. Hensley, T. Scheuermann, G. Coombe, M. Singh, A. Lastra, “Fast Summed-Area Table Generation and its Applications”, Computer Graphics Forum, 2005
- ⌚ [Atty05] L. Atty, N. Holzschuch, M. Lapierre, J.-M. Hasenfratz, C. Hansen, F. Sillion, “Soft Shadow Maps: Efficient Sampling of Light Source Visibility”, Technical Report, 2005